

PROJECT CONFIGURATION

Generally:

The project will be implemented in **Apache Hadoop**, with **Map-Reduce** programming technique, in **JAVA** programming language and will be tested locally in **pseudo-distributed mode**.

There is no limitation as to which OS or which IDE you will use as long as everyone uses the same hadoop version 2.7.3 and is able to run the final program in pseudo-distributed mode. Instructions below are given for Windows and Linux Operating Systems with Eclipse Luna IDE.

Prerequisites:

- 1) HADOOP
[hadoop-2.7.3.tar.gz](#)
- 2) JAVA
[JDK 8](#)
- 3) IDE
[Eclipse Luna](#) (recommended)
[NetBeans](#)
[IntelliJ](#)

HADOOP SETUP

Linux Hadoop Setup:

Follow detailed instructions.

Windows Hadoop Setup:

1. Download [hadoop-2.7.3.tar.gz](#) , unzip it in C:\
2. Add environmental user and system variables (right click MyComputer->Properties->Change settings->Advanced->Environmental Variables)
 - a. User variables-> New
Variable name: HADOOP_HOME Variable value: C:\hadoop-2.7.3\bin
 - b. System Variables-> Path -> Edit->New
C:\hadoop-2.7.3\bin
3. Download [JDK 8](#), install it
4. Go to folder C:/Program Files/Java/jdk1.8.0_121 and make sure you have installed Java successfully
5. Add environmental system variables (right click MyComputer->Properties->Change settings->Advanced->Environmental Variables)
 - a. System Variables-> New
Variable name: JAVA_HOME Variable value: C:\Progra~1\Java\jdk1.8.0_121
 - b. System Variables-> Path -> Edit->New
%JAVA_HOME%\bin
6. Go to C:\hadoop-2.7.3\etc\hadoop edit core-site.xml

```
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```
7. Go to C:\hadoop-2.7.3\etc\hadoop make a copy of mapred-site.xml.template, rename it to mapred-site.xml and edit it with:

```
<configuration>
  <property>
    <name> mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

8. Go to C:\hadoop-2.7.3\etc\hadoop edit hdfs-site.xml

```
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.namenode.dir</name>
    <value>C:\hadoop-2.7.3\data\namenode</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>C:\hadoop-2.7.3\data\namenode</value>
  </property>
</configuration>
```

9. Go to C:\hadoop-2.7.3 make new folder data

10. Go to C:\hadoop-2.7.3\etc\hadoop edit yarn-site.xml

```
<configuration>
  <property>
    <name>yarn.nodemanager.aux-services</name>
    <value>mapreduce_shuffle</value>
  </property>
  <property>
    <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
    <value>org.apache.hadoop.mapredShuffleHandler</value>
  </property>
</configuration>
```

11. Go to C:\hadoop-2.7.3\etc\hadoop edit hadoop-env.cmd

Change line 25 to: set JAVA_HOME= C:\Progra~1\Java\jdk1.8.0_121

12. Download [zip file](#), unzip it, copy files contained in hadooponwindows-master/bin folder and paste them in C:\hadoop-2.7.3\bin folder (replace files in the destination folder)

13. Open cmd and execute:

- a. `hadoop -version`
- b. `javac -version (1.8.0_121)`
- c. `java -version ("1.8.0_121")`
- d. `hdfs namenode -format`
- e. go to C:\hadoop-2.7.3\sbin>
 - i. `start-all.cmd` (should open 4 windows-deamons)
 - ii. `jps`

14. Open browser, 2 tabs

- a. `localhost:8088`
- b. `localhost:50070`

HADOOP is now running..

To shut down HADOOP type:

C:\hadoop-2.7.3\sbin>`stop-all.cmd`
(should stop 4 windows-deamons)

Hadoop Integration with Eclipse (Optional, Recommended)

1. Download [hadoop-eclipse-plugin-2.6.0.jar](#)
2. Copy the Map-Reduce plugin for eclipse in the plugins directory of your eclipse folder
C:\eclipse\eclipse\plugins (windows)
sudo cp /home/hdfs/Downloads/hadoop-eclipse-plugin-2.6.0.jar /opt/eclipse/plugins/ (linux)
3. Restart the eclipse using the command – /opt/eclipse/eclipse -vm
/usr/local/jdk1.8.0_121/bin/java -vmargs -Xmx1024m If eclipse is not coming up because of the X11 forwarding issue, try using “sux” instead of “su” while switching to the “hdfs”. (linux)
4. Start the eclipse: \$ECLIPSE_HOME/eclipse (linux)
5. In Eclipse menu click, Window --> Open Perspective --> Others --> MapReduce (both)
6. In bottom MapReduce icon click to Add new Hadoop location (both)
7. Enter MapReduce & HDFS running port (both)
Map/Reduce(V2) Master Host: localhost Port: 10020
DFS Master: localhost Port: 9000
8. Once Hadoop location added, DFS Locations will be seen/displayed in Eclipse Project Explorer window, (Windows-->Show View-->Project Explorer) (both)
9. Once Hadoop added, DFS Locations will be seen/displayed in Project Explorer window,
10. Right click DFS location and click to Connect (both)
11. Once connected successfully, it will display all the DFS Folder. (both)
12. Create folder named input. Create folder named output. Create a text file named WordCountSample.txt (any English content will do) upload the file in the input folder.

Running the WordCount Example with Eclipse and Maven (Optional, Recommended)

Configure Maven:

1. Download [apache-maven-3.3.9-bin.zip](#) unzip in C:\
2. Add environmental user and system variables (right click MyComputer->Properties->Change settings->Advanced->Environmental Variables)
 - a. User variables-> New
Variable name: MAVEN_HOME Variable value C:\apache-maven-3.3.9\bin
 - b. System Variables-> Path -> Edit->New
%MAVEN_HOME%\bin
3. Open cmd prompt as administrator and type below command to verify if maven installation.
 - a. C:\Windows\System32>mvn -version
Apache Maven 3.3.9 (bb52d8502b132ec0a5a3f4c09453c07478323dc5; 2015-11-10T18:41:47+02:00)
Maven home: C:\apache-maven-3.3.9\bin\..
Java version: 1.8.0_121, vendor: Oracle Corporation
Java home: C:\Program Files\Java\jdk1.8.0_121\jre
Default locale: en_US, platform encoding: Cp1253
OS name: "windows 10", version: "10.0", arch: "amd64", family: "dos"

Configuring New Maven Project in Eclipse:

1. Start Eclipse, go to Window->Open Perspective->Other-> Java and click ok
2. Right Click on the package explorer and select New->Other->Maven Project
3. Click next and check Use default workspace location and next
4. Select maven-archetype-quickstart and next
5. Specify archetype parameters
 - a. Group Id: WordCount
 - b. Artifact Id: HadoopMapReduce
 - c. Package: tuc.softnet.hadoop.mapreduce.example

6. Copy paste to the pom file the following:

```
<project xmlns="http://maven.apache.org/POM/4.0.0" xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
xsi:schemaLocation="http://maven.apache.org/POM/4.0.0 http://maven.apache.org/xsd/maven-4.0.0.xsd">
<modelVersion>4.0.0</modelVersion>

<groupId>WordCount</groupId>
<artifactId>HadoopMapReduce</artifactId>
<version>0.0.1-SNAPSHOT</version>
<packaging>jar</packaging>
<name>HadoopMapReduce</name>
<url>http://maven.apache.org</url>
<properties>
<project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
</properties>
<dependencies>
<dependency>
<groupId>junit</groupId>
<artifactId>junit</artifactId>
<version>3.8.1</version>
<scope>test</scope>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-annotations</artifactId>
<version>2.7.3</version>
</dependency>
<!-- Hadoop MapReduce Client Core -->
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-mapreduce-client-core</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-common</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-hdfs</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-hdfs</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>junit</groupId>
<artifactId>junit</artifactId>
<version>3.8.1</version>
<scope>test</scope>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-common</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-mapreduce-client-common</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-mapreduce-client-app</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-mapreduce-client-hs</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-mapreduce-client-hs-plugins</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-api</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-server-web-proxy</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-server-sharedcachemanager</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-server-resourcemanager</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-server-nodemanager</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-server-common</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-server-applicationhistoryservice</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-registry</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-common</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-client</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-applications-unmanaged-am-launcher</artifactId>
<version>2.7.3</version>
</dependency>
<dependency>
<groupId>org.apache.hadoop</groupId>
<artifactId>hadoop-yarn-applications-distributedshell</artifactId>
<version>2.7.3</version>
</dependency>
</dependencies>
</project>
```

7. Create class WordCountMapper.java

```
package tuc.softnet.hadoop.mapreduce.example;

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;

public class WordCountMapper extends Mapper<LongWritable, Text, Text, IntWritable> {

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> collector, Reporter reporter)
        throws IOException {
        String line = value.toString();
        StringTokenizer st = new StringTokenizer(line, " ");
        while (st.hasMoreTokens()) {
            word.set(st.nextToken());
            collector.collect(word, one);
        }
    }
}
```

8. Create class WordCountReducer.java

```
package tuc.softnet.hadoop.mapreduce.example;

import java.io.IOException;
import java.util.Iterator;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.Mapper;

public class WordCountReducer extends Mapper<Text, IntWritable, Text, IntWritable> {

    public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> outputCollector,
        Reporter reporter) throws IOException {
        int sum = 0;

        while (values.hasNext()) {
            sum = sum + values.next().get();
        }
        outputCollector.collect(key, new IntWritable(sum));
    }
}
```

9. Create class WordCount.java

```
package tuc.softnet.hadoop.mapreduce.example;

import java.net.URI;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.RunningJob;
import org.apache.hadoop.mapred.TextInputFormat;

public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();

        Path inputPath = new Path("hdfs://127.0.0.1:9000/input/WordCountSample.txt");
        Path outputPath = new Path("hdfs://127.0.0.1:9000/output/result");

        JobConf job = new JobConf(conf, WordCount.class);
        job.setJarByClass(WordCount.class);
        job.setJobName("WordCounterJob");

        FileInputFormat.setInputPaths(job, inputPath);
        FileOutputFormat.setOutputPath(job, outputPath);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setOutputFormat(TextOutputFormat.class);
        job.setMapperClass(WordCountMapper.class);
        job.setReducerClass(WordCountReducer.class);

        FileSystem hdfs = FileSystem.get(URI.create("hdfs://127.0.0.1:9000/"), conf);
        if (hdfs.exists(outputPath))
            hdfs.delete(outputPath, true);

        RunningJob runningJob = JobClient.runJob(job);
        System.out.println("job isSuccessful: " + runningJob.isComplete());
    }
}
```

10. Right click on WordCount.java -> Run As -> Run on Hadoop

log4j:WARN No appenders could be found for logger (org.apache.hadoop.metrics2.lib.MutableMetricsFactory).

log4j:WARN Please initialize the log4j system properly.

log4j:WARN See <http://logging.apache.org/log4j/1.2/faq.html#noconfig> for more info.

job.isSuccessful: true

11. Right click on DFS Location and Disconnect and Refresh, you should see output file generated in localhost->output->result

OTHER USEFUL LINKS

1. [HADOOP tutorial](#) for linux
2. [HDFS commands](#) shell or command line.